

You are working as a machine learning engineer to develop a neural network-based classifier for detecting cars in grayscale images. Each image is  $50 \times 50$  pixels, represented as a 2,500-dimensional feature vector. You will follow the steps of forward propagation, backpropagation, and evaluation, performing key calculations at each stage.

### Task 1: Forward Propagation (Computation Required)

#### Given:

- Input feature vector  $X$  of size  $2500 \times 1$ .
- Weight matrix  $W^{(1)}$  for the first hidden layer has dimensions  $128 \times 2500$ .
- Bias vector  $b^{(1)}$  is of size  $128 \times 1$ .
- Activation function in the hidden layer: ReLU ( $\text{ReLU}(z) = \max(0, z)$ ).
- Second layer has 1 neuron with a sigmoid activation.
- Weight matrix  $W^{(2)}$  has size  $1 \times 128$ .
- Bias  $b^{(2)}$  is a scalar.

#### Your Tasks:

1. Compute the first layer activation  $A^{(1)}$  using:

$$Z^{(1)} = W^{(1)}X + b^{(1)}$$

$$A^{(1)} = \max(0, Z^{(1)})$$

Assume:

$$W^{(1)} = \begin{bmatrix} 0.02 & -0.01 & \dots & 0.005 \\ -0.03 & 0.04 & \dots & -0.002 \\ \vdots & \vdots & \ddots & \vdots \\ 0.01 & -0.02 & \dots & 0.03 \end{bmatrix},$$

$b^{(1)} = \mathbf{0}$  (vector of zeros).

2. Compute the second layer output:

$$Z^{(2)} = W^{(2)}A^{(1)} + b^{(2)}$$

Assume:

$$W^{(2)} = [0.05 \quad -0.02 \quad 0.01 \quad \dots \quad 0.02],$$

$b^{(2)} = -0.1$ .

3. Apply **sigmoid activation** to obtain the final prediction  $A^{(2)}$ :

$$A^{(2)} = \frac{1}{1 + e^{-Z^{(2)}}}$$

Compute the **final predicted probability**  $A^{(2)}$  given a random input  $X$  where each feature is drawn from a uniform distribution between 0 and 1.

### Task 2: Backpropagation (Compute Gradients)

Using the **cross-entropy loss function**:

$$J = -[y \log(A^{(2)}) + (1 - y) \log(1 - A^{(2)})]$$

where  $y = 1$  if the image contains a car and  $y = 0$  otherwise.

**Your Tasks:**

1. Compute the derivative of the loss with respect to  $Z^{(2)}$ :

$$\frac{\partial J}{\partial Z^{(2)}} = A^{(2)} - y$$

Assume  $y = 1$ .

2. Compute the gradient of the cost function with respect to  $W^{(2)}$ :
3. Compute the gradient of the cost function with respect to  $W^{(1)}$  using **backpropagation through ReLU**:

$$\frac{\partial J}{\partial W^{(1)}} = \left( \left( W^{(2)} \right)^T \frac{\partial J}{\partial Z^{(2)}} \right) \circ 1(Z^{(1)} > 0) X^T$$

where  $\circ$  represents element-wise multiplication.

Compute the gradients for  $W^{(2)}$  and  $W^{(1)}$  based on the forward propagation output from Task 1. And tell me the shape of dW2, minimum element of dW2 and maximum element of dB2.

**Task 3: Model Evaluation**

You trained the neural network for 10 epochs and obtained the following confusion matrix on the test set:

	<b>Predicted: Car (1)</b>	<b>Predicted: Not Car (0)</b>
Actual: Car (1)	150	50
Actual: Not Car (0)	30	270

Calculate the accuracy, precision, and recall of the classifier using the confusion matrix.

**Task 4: Multi-class Classification**

Your task is now to extend the model to classify **four different vehicle types**:

- Pedestrian (Class 1)
- Car (Class 2)
- Motorcycle (Class 3)
- Truck (Class 4)

Instead of a single output neuron, you now use a softmax output layer with 4 neurons.

Your task:

- Implement the **softmax function** for output layer activation:

$$A_j = \frac{e^{Z_j}}{\sum_{k=1}^4 e^{Z_k}}, \quad \forall j \in \{1, 2, 3, 4\}$$

Given:

$$Z = [2.1, 1.4, 0.5, -0.2]$$

2. Ensure numerical stability by subtracting the **maximum value** from all logits:

$$Z' = Z - \max(Z)$$

Compute the probability of each class using the softmax activation function.

### General instruction for your Python Code:

```
import numpy as np

def sigmoid(z):
    #YOUR CODE HERE

def relu(z):
    #YOUR CODE HERE

def softmax(z):
    #YOUR CODE HERE

def forward_propagation(X, W1, b1, W2, b2):
    #YOUR CODE HERE

def compute_gradients(X, A1, A2, Y, W2, Z1):
    #YOUR CODE HERE

def evaluate_model(confusion_matrix):
    #YOUR CODE HERE

# Fixed Data (Non-Random)
X = np.ones((2500, 1)) * 0.5 # Fixed input values
W1 = np.linspace(-0.01, 0.01, num=2500 * 128).reshape(128, 2500)
b1 = np.zeros((128, 1))
W2 = np.linspace(-0.01, 0.01, num=128).reshape(1, 128)
b2 = np.zeros((1, 1))
Y = np.array([[1]]) # Assume car is present

# Forward propagation
Z1, A1, Z2, A2 = forward_propagation(X, W1, b1, W2, b2)
print("Predicted Probability (A2):", A2)

# Compute gradients
dW1, dB1, dW2, dB2 = compute_gradients(X, A1, A2, Y, W2, Z1)
print("Gradient dW2 shape:", dW2.shape)

# Confusion Matrix
conf_matrix = np.array([[150, 50], [30, 270]])
accuracy, precision, recall = evaluate_model(conf_matrix)
print(f"Accuracy: {accuracy:.2f}, Precision: {precision:.2f}, Recall: {recall:.2f}")

# Multi-Class Softmax Example
Z_multiclass = np.array([2.1, 1.4, 0.5, -0.2]).reshape(-1, 1)
A_softmax = softmax(Z_multiclass)
print("Softmax Probabilities:", A_softmax.flatten())
```

import numpy as np

```

def sigmoid(z):
    #YOUR CODE HERE

def relu(z):
    #YOUR CODE HERE

def softmax(z):
    #YOUR CODE HERE

def forward_propagation(X, W1, b1, W2, b2):
    #YOUR CODE HERE

def compute_gradients(X, A1, A2, Y, W2, Z1):
    #YOUR CODE HERE

def evaluate_model(confusion_matrix):
    #YOUR CODE HERE


# Fixed Data (Non-Random)
X = np.ones((2500, 1)) * 0.5 # Fixed input values
W1 = np.linspace(-0.01, 0.01, num=2500 * 128).reshape(128, 2500)
b1 = np.zeros((128, 1))
W2 = np.linspace(-0.01, 0.01, num=128).reshape(1, 128)
b2 = np.zeros((1, 1))
Y = np.array([[1]]) # Assume car is present


# Forward propagation
Z1, A1, Z2, A2 = forward_propagation(X, W1, b1, W2, b2)
print("Predicted Probability (A2):", A2)


# Compute gradients
dW1, dB1, dW2, dB2 = compute_gradients(X, A1, A2, Y, W2, Z1)
print("Gradient dW2 shape:", dW2.shape)


# Confusion Matrix
conf_matrix = np.array([[150, 50], [30, 270]])
accuracy, precision, recall = evaluate_model(conf_matrix)
print(f'Accuracy: {accuracy:.2f}, Precision: {precision:.2f}, Recall: {recall:.2f}')


# Multi-Class Softmax Example
Z_multiclass = np.array([2.1, 1.4, 0.5, -0.2]).reshape(-1, 1)
A_softmax = softmax(Z_multiclass)
print("Softmax Probabilities:", A_softmax.flatten())

```